

Project_NeuralNetwork_implementation-Copy1

July 1, 2025

```
[1]: %%capture
!pip install tensorflow
!pip install keras_tuner
```

```
[2]: %%capture

import pandas as pd
import time
import numpy as np
import os
import requests
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.callbacks import EarlyStopping
import keras_tuner as kt
```

```
2025-07-01 17:57:29.613960: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-07-01 17:57:29.620680: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-07-01 17:57:29.636242: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1751392649.662243      221 cuda_dnn.cc:8579] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1751392649.669664      221 cuda_blas.cc:1407] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
W0000 00:00:1751392649.689751      221 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1751392649.689785      221 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
```

```
target more than once.
W0000 00:00:1751392649.689788      221 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
W0000 00:00:1751392649.689790      221 computation_placer.cc:177] computation
placer already registered. Please check linkage and avoid linking the same
target more than once.
2025-07-01 17:57:29.696434: I tensorflow/core/platform/cpu_feature_guard.cc:210]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
```

1 Original Data Source

The original data for this project is sourced from <https://optiondata.org/>. They have graciously provided free sample datasets for option prices and stock prices from January 2013 to June 2013. The data we will be using is from the month of February 2013 located in the 2013-02.zip file located on their website.

1.1 Data Preprocessing

The dataset we will be using is fairly large, there are 19 trading days in the month of February 2013 and two different csv files for each day, the first being the options data and the second being the stock price data.

To simplify this, we will first join the two datasets through the `symbol` column in the stocks files, and the `underlying` column in the options files. This will allow us to access the data for both in a single file for our features later.

After reviewing the files, we found there were around 500,000 rows for 3800 unique tickers on each trading day, with 19 different days that meant we were looking at close to 9.5 million rows of data. Training machine learning algorithms on this much data will take a very long time, so to minimize this effect, we decided to filter the data to a subset of S&P 100 stocks as of February 1st 2013 and filter the options to call options only, allowing us to train our models within a reasonable time-frame.

Lastly, since our original options data set contained bid and ask prices, we engineered a new feature named `mid_price` which contains the average of the two, and will be evaluated as the price of the option.

Below I've attached a code chunk which displays the filtering we performed. (Please note this won't work on your machine unless you've modified the directories with your own):

```
import numpy as np
import pandas as pd
import os

def combine_options_data(options_data, stock_data):
    options_data['mid_price'] = (options_data['bid'] + options_data['ask'])/2
```

```

combined_df = options_data.merge(stock_data,
                                left_on = "underlying", right_on = "symbol",
                                suffixes=('_', '_stock'))
combined_df = combined_df.drop(columns=['symbol'])
return combined_df

def combine_new_data(date):
    options = pd.read_csv(f"/home/steve/Downloads/CFRM_521/ProjectData/2013-02/{date}options.csv")
    stocks = pd.read_csv(f"/home/steve/Downloads/CFRM_521/ProjectData/2013-02/{date}stocks.csv")
    return combine_options_data(options, stocks)

url = "https://web.archive.org/web/20130201003232/https://en.wikipedia.org/wiki/S%26P_100"
sp100_comp = pd.read_html(url)
sp100_comp = sp100_comp[2]["Symbol"]
sp100_comp = sp100_comp.unique()

dir = "/home/steve/Downloads/CFRM_521/ProjectData/2013-02/"
dates = []
for file in os.listdir(dir):
    #print(file)
    if file.endswith(".csv"):
        if "options" in file:
            dates.append(file.split("options")[0])

dates = sorted(set(dates))
for date in dates:
    print(f"Processing: {date}")
    combined_df = combine_new_data(date)
    combined_df['underlying'] = (
        combined_df['underlying']
        .str.replace('.', '-', regex=False)
        .str.upper()
        .replace({'GOOGL': 'GOOG'})
    )
    sp100_comp = [sym.replace('.', '-').upper() for sym in sp100_comp]
    filt_df = combined_df[combined_df['underlying'].isin(sp100_comp)]
    num_stocks = filt_df['underlying'].unique()
    if len(num_stocks) != len(sp100_comp):
        print(f"Error size mismatch, Filtered Df Size: {len(num_stocks)}, SP100 Size: {len(sp100_comp)}")
        missing = set(sp100_comp) - set(filt_df['underlying'].unique())
        print(f"Missing tickers: {sorted(missing)}")
    else:
        filt_df.to_csv(f"/home/steve/Downloads/CFRM_521/ProjectData/filtered/{date}.csv", index=False)

```

2 Loading in Data

The preprocessed data was then uploaded to a github repository to allow all team-members to access the new data. The code below is used to import the data into the jupyter notebook and then randomize the order in which the rows appear.

```
[3]: url = "https://api.github.com/repos/stevedemirev/CFRM521-ProjectData/contents/
      ↪filtered"
      response = requests.get(url)
      files = response.json()

      csv_files = sorted([file for file in files if file['name'].endswith('.csv')],
      ↪key = lambda x: x['name'])

      def get_datasets(files):
          df = pd.DataFrame()
          for file in files:
              temp = pd.read_csv(file['download_url'])
              df = pd.concat([df, temp], ignore_index = True)
          return df

      full_df = get_datasets(csv_files)
      full_df = full_df.sample(frac=1, random_state=42).reset_index(drop=True)

      total_files = len(full_df)
      train_size = int(total_files*0.7)
      val_size = int(total_files*0.85)

      train = full_df[:train_size]
      valid = full_df[train_size:val_size]
      test = full_df[val_size:]
```

```
[4]: display(train.head())
```

	contract	underlying	expiration	type	strike	style	bid	\
0	V130921C00110000	V	2013-09-21	call	110.0	A	48.10	
1	GOOG150117C00370000	GOOG	2015-01-17	call	370.0	A	415.00	
2	AMZN150117P00240000	AMZN	2015-01-17	put	240.0	A	36.15	
3	AAPL130720P00630000	AAPL	2013-07-20	put	630.0	A	189.45	
4	BAC130322C00013500	BAC	2013-03-22	call	13.5	A	0.00	

	bid_size	ask	ask_size	...	theta	vega	implied_volatility	\
0	NaN	48.75	NaN	...	-1.5975	6.2814	0.2444	
1	NaN	419.40	NaN	...	-4.2084	43.9866	0.2792	
2	NaN	37.05	NaN	...	-11.5955	130.5729	0.3374	
3	NaN	190.40	NaN	...	-11.9728	21.8567	0.2993	
4	NaN	0.02	NaN	...	-0.5092	0.1719	0.3651	

	mid_price	open	high	low	close	volume_stock \
0	48.425	156.350006	158.080002	155.740005	157.990005	17884400
1	417.200	778.400030	783.000040	773.750022	782.419993	4331200
2	36.600	260.890015	262.040009	255.729996	259.359985	3348600
3	189.925	453.850014	455.120003	442.569996	442.799988	93144800
4	0.010	11.150000	11.360000	11.100000	11.300000	147145400

	adjust_close
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 25 columns]

```
[5]: display(train.tail())
```

	contract	underlying	expiration	type	strike	style	bid \
483143	COST140118P00100000	COST	2014-01-18	put	100.00	A	6.50
483144	MCD140118C00085000	MCD	2014-01-18	call	85.00	A	10.60
483145	EMR140118C00025000	EMR	2014-01-18	call	25.00	A	30.70
483146	GILD130817P00036250	GILD	2013-08-17	put	36.25	A	1.33
483147	GD140118C00040000	GD	2014-01-18	call	40.00	A	27.40

	bid_size	ask	ask_size	...	theta	vega	implied_volatility \
483143	NaN	6.70	NaN	...	-3.8113	38.6961	0.1933
483144	NaN	10.80	NaN	...	-1.0126	27.5612	0.1603
483145	NaN	33.70	NaN	...	-0.6678	5.9553	0.6433
483146	NaN	1.38	NaN	...	-2.5770	9.1628	0.2868
483147	NaN	29.40	NaN	...	-0.4518	2.2104	0.2716

	mid_price	open	high	low	close \
483143	6.600	102.879997	103.190002	100.940002	101.699997
483144	10.700	94.269997	95.250000	93.849998	95.250000
483145	32.200	57.330002	57.509998	56.840000	57.470001
483146	1.355	40.459999	40.950001	40.299999	40.830002
483147	28.400	68.449997	69.040001	67.959999	67.970001

	volume_stock	adjust_close
483143	3192000	NaN
483144	4421300	NaN
483145	2912700	NaN
483146	9791000	NaN
483147	2338200	NaN

[5 rows x 25 columns]

```
[6]: total_size = len(train) + len(valid) + len(test)
def proportion(df):
    return round(len(df)/total_size,3)

print(f"Length of Training set: {len(train):,} rows, Proportion:␣
↪{proportion(train)}")
print(f"Length of Validation set: {len(valid):,} rows, Proportion:␣
↪{proportion(valid)}")
print(f"Length of Testing set: {len(test):,} rows, Proportion:␣
↪{proportion(test)}")
print(f"Original Dataset size: {total_size:,} rows, Sum Check:␣
↪{proportion(train)+proportion(valid)+proportion(test)}")
```

```
Length of Training set: 483,148 rows, Proportion: 0.7
Length of Validation set: 103,532 rows, Proportion: 0.15
Length of Testing set: 103,532 rows, Proportion: 0.15
Original Dataset size: 690,212 rows, Sum Check: 1.0
```

3 Baseline Model

For my implementation, I will be using the `close`, `strike`, `delta`, `gamma`, `vega`, `theta`, and `implied_volatility` columns as features. I also engineered `tte` (time to expiry) as it's an important parameter used in one the most common methods of pricing options, the Black-Scholes formula.

For my model's architecture, I decided to use 4 hidden layers with 50 neurons each with 'ReLU' activation, 'he_normal' initializer, 'Nadam' optimizer, early stopping with a patience of 10, and "MSE" as the primary loss metric, and "MAE" as the secondary loss metric. The choice to use 4 hidden layers with 50 neurons was somewhat arbitrary and not backed by any theoretical justification other than the expectation of it being able to accurately capture any non-linear relationships occurring with options pricing. The output layer would remain the default linear activation function with one neuron as we are expecting a single value, the mid price. I also chose the "ReLU" activation function for each neuron with "he_normal" initializer to better model the non-linear relationship and avoid vanishing gradients during training. Additionally, "ReLU" works well because it encourages positive only outputs, which matches with the intrinsic value of an option is defined as: $V = \max(S_T - K, 0)$. The 'Nadam' optimizer was chosen here as its learning rate is adaptable, allowing it to converge faster and smoother for non-linear functions compared to other optimizers. Early stopping was added as a form of regularization to prevent unnecessary training once the model stops improving on the validation set. Given the time required to train the model, early stopping is an effective way to reduce training time while also helping to avoid overfitting. Since the objective of this project is to predict option prices, Mean Squared Error (MSE) and Mean Absolute Error (MAE) are appropriate evaluation metrics. MSE penalizes larger errors more heavily, making it useful for identifying significant prediction deviations, while MAE provides a more interpretable measure of the average prediction error.

```
[7]: def define_features(df):
    df['expiration'] = pd.to_datetime(df['expiration'])
```

```

df['quote_date'] = pd.to_datetime(df['quote_date'])
df['tte'] = (df['expiration'] - df['quote_date']).dt.days / 252
X = df[['close', 'strike', 'delta', 'gamma',
        'vega', 'theta', 'implied_volatility', 'tte']]
y = df['mid_price']
return X, y

def get_features(opt_type):
    X_feats = []
    y_feats = []
    for df in [train, valid, test]:
        df_temp = df[df['type'] == opt_type].copy()
        X, y = define_features(df_temp)
        X_feats.append(X)
        y_feats.append(y)
    return X_feats, y_feats

# Calls
X_feats, y_feats = get_features("call")
X_train_c, X_valid_c, X_test_c = X_feats
y_train_c, y_valid_c, y_test_c = y_feats

# Puts
X_feats, y_feats = get_features("put")
X_train_p, X_valid_p, X_test_p = X_feats
y_train_p, y_valid_p, y_test_p = y_feats

```

```

[8]: def scale_data(train, val, test):
    scaler = StandardScaler()
    train_scaled = scaler.fit_transform(train)
    valid_scaled = scaler.transform(val)
    test_scaled = scaler.transform(test)
    return train_scaled, valid_scaled, test_scaled

def reset_session(seed=42):
    tf.keras.backend.clear_session()
    tf.random.set_seed(seed)
    np.random.seed(seed)

def build_model(input_shape):
    reset_session()
    model = tf.keras.Sequential([
        tf.keras.Input(shape = (input_shape,)),
        tf.keras.layers.Dense(50, activation = "relu", kernel_initializer =_
↪ "he_normal"),
        tf.keras.layers.Dense(50, activation = "relu", kernel_initializer =_
↪ "he_normal"),

```

```

        tf.keras.layers.Dense(50, activation = "relu", kernel_initializer =
↪ "he_normal"),
        tf.keras.layers.Dense(50, activation = "relu", kernel_initializer =
↪ "he_normal"),
        tf.keras.layers.Dense(1)
    ])

    model.compile(
        optimizer = "nadam",
        loss = "mse",
        metrics = ['mae']
    )
    return model

X_train_scaled_c, X_valid_scaled_c, X_test_scaled_c = scale_data(X_train_c,
↪ X_valid_c, X_test_c)
X_train_scaled_p, X_valid_scaled_p, X_test_scaled_p = scale_data(X_train_p,
↪ X_valid_p, X_test_p)

early_stop = EarlyStopping(
    monitor = "val_loss",
    patience = 10,
    mode = "min",
    restore_best_weights = True
)

```

3.1 Call Option Model

```

[9]: model_calls = build_model(X_train_scaled_c.shape[1])

start = time.time()
history_calls = model_calls.fit(X_train_scaled_c, y_train_c,
                                validation_data = (X_valid_scaled_c, y_valid_c),
                                epochs = 50, verbose = 1, callbacks=[early_stop])
end = time.time()
baseline_time_c = end-start

mse_val_c, mae_val_c = model_calls.evaluate(X_valid_scaled_c,
                                             y_valid_c, verbose = 0)
test_loss_c, test_mae_c = model_calls.evaluate(X_test_scaled_c,
                                                y_test_c, verbose = 0)

print("\nCall Options:")
print(f"Validation MSE: {mse_val_c}")
print(f"Validation MAE: {mae_val_c}")

```



```
print(f"\nTest MSE: {test_loss_c}")
print(f"Test MAE: {test_mae_c}")
```

Epoch 1/50

2025-07-01 17:57:51.221715: E

external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:51] failed call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: UNKNOWN ERROR (303)

7549/7549 22s 3ms/step -
loss: 276.4209 - mae: 4.1048 - val_loss: 2.5094 - val_mae: 0.8404

Epoch 2/50

7549/7549 19s 2ms/step -
loss: 2.6091 - mae: 0.8065 - val_loss: 1.7289 - val_mae: 0.6684

Epoch 3/50

7549/7549 18s 2ms/step -
loss: 1.8883 - mae: 0.6353 - val_loss: 1.5641 - val_mae: 0.5589

Epoch 4/50

7549/7549 19s 2ms/step -
loss: 1.6471 - mae: 0.5632 - val_loss: 1.5665 - val_mae: 0.5549

Epoch 5/50

7549/7549 18s 2ms/step -
loss: 1.5697 - mae: 0.5330 - val_loss: 1.1716 - val_mae: 0.4646

Epoch 6/50

7549/7549 19s 2ms/step -
loss: 1.4368 - mae: 0.4939 - val_loss: 1.1193 - val_mae: 0.4454

Epoch 7/50

7549/7549 19s 3ms/step -
loss: 1.3244 - mae: 0.4629 - val_loss: 1.0621 - val_mae: 0.4049

Epoch 8/50

7549/7549 21s 3ms/step -
loss: 1.3028 - mae: 0.4508 - val_loss: 0.8277 - val_mae: 0.3295

Epoch 9/50

7549/7549 20s 3ms/step -
loss: 1.2889 - mae: 0.4364 - val_loss: 1.2629 - val_mae: 0.4374

Epoch 10/50

7549/7549 22s 3ms/step -
loss: 1.2268 - mae: 0.4206 - val_loss: 0.9297 - val_mae: 0.3587

Epoch 11/50

7549/7549 19s 3ms/step -
loss: 1.1992 - mae: 0.4104 - val_loss: 1.1181 - val_mae: 0.3986

Epoch 12/50

7549/7549 20s 3ms/step -
loss: 1.1733 - mae: 0.3992 - val_loss: 1.0529 - val_mae: 0.3779

Epoch 13/50

7549/7549 19s 3ms/step -
loss: 1.1645 - mae: 0.3946 - val_loss: 0.9847 - val_mae: 0.3546

Epoch 14/50

7549/7549 18s 2ms/step -
 loss: 1.1477 - mae: 0.3882 - val_loss: 1.0925 - val_mae: 0.3883
 Epoch 15/50
 7549/7549 18s 2ms/step -
 loss: 1.1308 - mae: 0.3830 - val_loss: 0.9918 - val_mae: 0.3623
 Epoch 16/50
 7549/7549 19s 2ms/step -
 loss: 1.1146 - mae: 0.3762 - val_loss: 0.7738 - val_mae: 0.2967
 Epoch 17/50
 7549/7549 19s 2ms/step -
 loss: 1.0782 - mae: 0.3666 - val_loss: 0.8522 - val_mae: 0.3380
 Epoch 18/50
 7549/7549 18s 2ms/step -
 loss: 1.0807 - mae: 0.3669 - val_loss: 0.7166 - val_mae: 0.2859
 Epoch 19/50
 7549/7549 18s 2ms/step -
 loss: 1.0669 - mae: 0.3602 - val_loss: 0.6996 - val_mae: 0.2739
 Epoch 20/50
 7549/7549 18s 2ms/step -
 loss: 1.0474 - mae: 0.3564 - val_loss: 0.6927 - val_mae: 0.2743
 Epoch 21/50
 7549/7549 18s 2ms/step -
 loss: 1.0251 - mae: 0.3493 - val_loss: 0.7833 - val_mae: 0.3282
 Epoch 22/50
 7549/7549 18s 2ms/step -
 loss: 1.0163 - mae: 0.3480 - val_loss: 0.6972 - val_mae: 0.2742
 Epoch 23/50
 7549/7549 18s 2ms/step -
 loss: 1.0194 - mae: 0.3477 - val_loss: 0.6800 - val_mae: 0.2694
 Epoch 24/50
 7549/7549 18s 2ms/step -
 loss: 1.0018 - mae: 0.3427 - val_loss: 0.7904 - val_mae: 0.3294
 Epoch 25/50
 7549/7549 18s 2ms/step -
 loss: 0.9894 - mae: 0.3389 - val_loss: 0.6555 - val_mae: 0.2678
 Epoch 26/50
 7549/7549 19s 3ms/step -
 loss: 0.9811 - mae: 0.3374 - val_loss: 0.7301 - val_mae: 0.2992
 Epoch 27/50
 7549/7549 21s 3ms/step -
 loss: 0.9760 - mae: 0.3349 - val_loss: 0.7146 - val_mae: 0.2959
 Epoch 28/50
 7549/7549 20s 3ms/step -
 loss: 0.9647 - mae: 0.3320 - val_loss: 0.7400 - val_mae: 0.3099
 Epoch 29/50
 7549/7549 19s 3ms/step -
 loss: 0.9550 - mae: 0.3287 - val_loss: 0.7331 - val_mae: 0.3103
 Epoch 30/50

```

7549/7549          19s 3ms/step -
loss: 0.9483 - mae: 0.3252 - val_loss: 0.6975 - val_mae: 0.2908
Epoch 31/50
7549/7549          19s 3ms/step -
loss: 0.9366 - mae: 0.3220 - val_loss: 0.7476 - val_mae: 0.3209
Epoch 32/50
7549/7549          18s 2ms/step -
loss: 0.9357 - mae: 0.3210 - val_loss: 0.7601 - val_mae: 0.3289
Epoch 33/50
7549/7549          20s 3ms/step -
loss: 0.9309 - mae: 0.3204 - val_loss: 0.6711 - val_mae: 0.2786
Epoch 34/50
7549/7549          19s 3ms/step -
loss: 0.9311 - mae: 0.3188 - val_loss: 0.6766 - val_mae: 0.2764
Epoch 35/50
7549/7549          20s 3ms/step -
loss: 0.9290 - mae: 0.3192 - val_loss: 0.6760 - val_mae: 0.2784

```

Call Options:

Validation MSE: 0.6554962396621704

Validation MAE: 0.2678316831588745

Test MSE: 0.8441799283027649

Test MAE: 0.26972681283950806

```

[10]: def plot_learning_curves(history, opt_type):
        fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(8, 6))

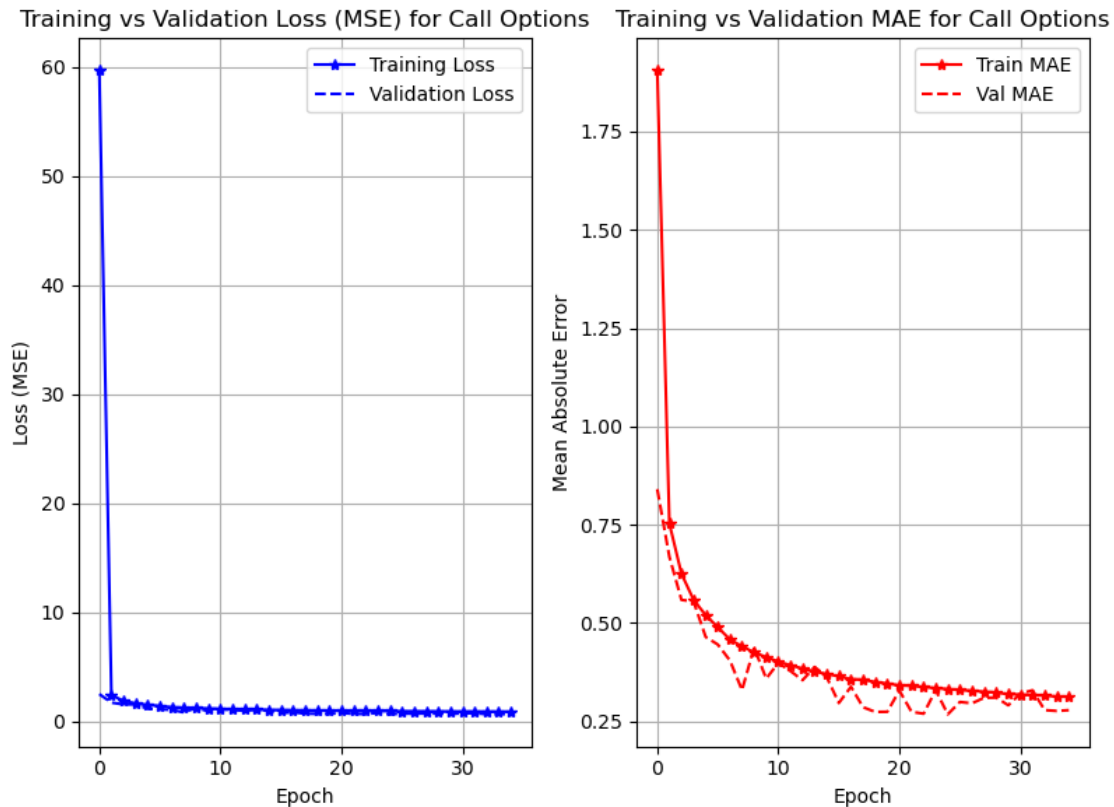
        ax1.plot(history.epoch, history.history["loss"], '*', label="Training Loss", color='b')
        ax1.plot(history.epoch, history.history["val_loss"], '--', label="Validation Loss", color='b')
        ax1.set_title(f"Training vs Validation Loss (MSE) for {opt_type} Options")
        ax1.set_xlabel("Epoch")
        ax1.set_ylabel("Loss (MSE)")
        ax1.grid(True)
        ax1.legend()

        ax2.plot(history.epoch, history.history["mae"], '*', label="Train MAE", color='r')
        ax2.plot(history.epoch, history.history["val_mae"], '--', label="Val MAE", color='r')
        ax2.set_title(f"Training vs Validation MAE for {opt_type} Options")
        ax2.set_xlabel("Epoch")
        ax2.set_ylabel("Mean Absolute Error")
        ax2.grid(True)
        ax2.legend()

```

```
plt.tight_layout()
plt.show()

plot_learning_curves(history_calls, "Call")
```



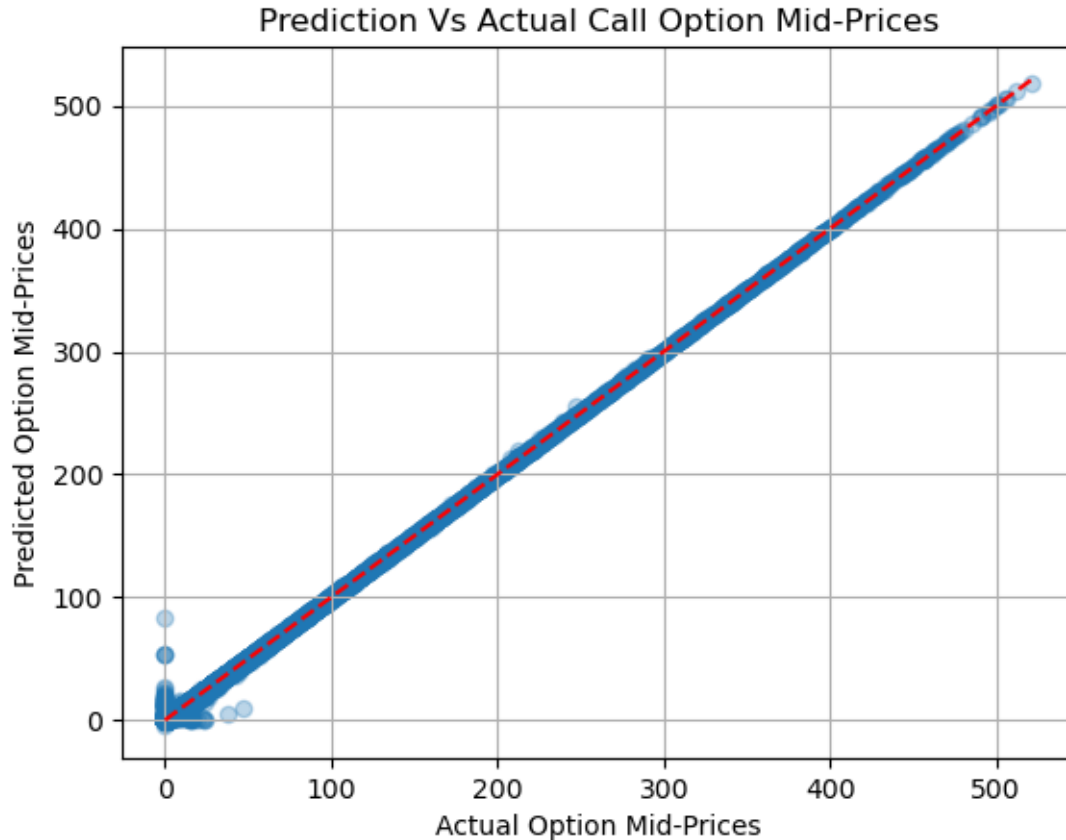
```
[11]: def get_prediction_plot(model, X_test_s, y_test, opt_type):
    y_pred = model.predict(X_test_s)

    plt.scatter(y_test, y_pred, marker = 'o', alpha = 0.3)
    plt.plot([y_test.min(), y_test.max()],
             [y_test.min(), y_test.max()], 'r--')
    plt.xlabel("Actual Option Mid-Prices")
    plt.ylabel("Predicted Option Mid-Prices")
    plt.title(f"Prediction Vs Actual {opt_type} Option Mid-Prices")
    plt.grid(True)
    plt.show()

get_prediction_plot(model_calls, X_test_scaled_c,
                    y_test_c, "Call")
```

1617/1617

2s 1ms/step



3.2 Put option Model

```
[12]: model_puts = build_model(X_train_scaled_p.shape[1])

start = time.time()
history_puts = model_puts.fit(X_train_scaled_p, y_train_p,
                             validation_data = (X_valid_scaled_p, y_valid_p),
                             epochs = 50, verbose = 1, callbacks=[early_stop])
end = time.time()
baseline_time_p = end-start

mse_val_p, mae_val_p = model_puts.evaluate(X_valid_scaled_p,
                                           y_valid_p, verbose = 0)
test_loss_p, test_mae_p = model_puts.evaluate(X_test_scaled_p,
                                              y_test_p, verbose = 0)

print("\nPut Options:")
print(f"Validation MSE: {mse_val_p}")
print(f"Validation MAE: {mae_val_p}")
```

```
print(f"\nTest MSE: {test_loss_p}")
print(f"Test MAE: {test_mae_p}")
```

```
Epoch 1/50
7550/7550          22s 3ms/step -
loss: 386.3810 - mae: 4.3048 - val_loss: 44.0751 - val_mae: 2.6941
Epoch 2/50
7550/7550          19s 2ms/step -
loss: 9.3228 - mae: 1.2542 - val_loss: 5.8460 - val_mae: 1.0468
Epoch 3/50
7550/7550          19s 3ms/step -
loss: 6.2965 - mae: 0.9776 - val_loss: 5.2344 - val_mae: 0.9160
Epoch 4/50
7550/7550          19s 2ms/step -
loss: 5.4064 - mae: 0.8472 - val_loss: 5.7089 - val_mae: 0.9302
Epoch 5/50
7550/7550          19s 3ms/step -
loss: 4.7906 - mae: 0.7692 - val_loss: 7.0501 - val_mae: 0.9922
Epoch 6/50
7550/7550          20s 3ms/step -
loss: 4.6544 - mae: 0.7236 - val_loss: 6.4352 - val_mae: 0.9558
Epoch 7/50
7550/7550          18s 2ms/step -
loss: 4.2419 - mae: 0.6785 - val_loss: 7.5897 - val_mae: 0.9951
Epoch 8/50
7550/7550          18s 2ms/step -
loss: 4.0587 - mae: 0.6474 - val_loss: 5.1543 - val_mae: 0.8332
Epoch 9/50
7550/7550          18s 2ms/step -
loss: 3.9256 - mae: 0.6358 - val_loss: 7.1388 - val_mae: 0.9846
Epoch 10/50
7550/7550          18s 2ms/step -
loss: 3.7227 - mae: 0.6076 - val_loss: 4.2217 - val_mae: 0.7353
Epoch 11/50
7550/7550          18s 2ms/step -
loss: 3.6603 - mae: 0.5964 - val_loss: 3.5941 - val_mae: 0.6683
Epoch 12/50
7550/7550          19s 2ms/step -
loss: 3.6692 - mae: 0.5867 - val_loss: 3.6855 - val_mae: 0.6136
Epoch 13/50
7550/7550          20s 3ms/step -
loss: 3.4678 - mae: 0.5648 - val_loss: 2.8578 - val_mae: 0.5301
Epoch 14/50
7550/7550          19s 2ms/step -
loss: 3.1924 - mae: 0.5359 - val_loss: 3.0826 - val_mae: 0.5284
Epoch 15/50
```

7550/7550 19s 3ms/step -
 loss: 3.2803 - mae: 0.5406 - val_loss: 2.8242 - val_mae: 0.5444
 Epoch 16/50
 7550/7550 20s 3ms/step -
 loss: 3.0732 - mae: 0.5232 - val_loss: 2.4605 - val_mae: 0.4808
 Epoch 17/50
 7550/7550 20s 3ms/step -
 loss: 3.0614 - mae: 0.5224 - val_loss: 2.4961 - val_mae: 0.4985
 Epoch 18/50
 7550/7550 20s 3ms/step -
 loss: 2.8887 - mae: 0.5019 - val_loss: 2.3405 - val_mae: 0.4731
 Epoch 19/50
 7550/7550 20s 3ms/step -
 loss: 2.8594 - mae: 0.4972 - val_loss: 2.3819 - val_mae: 0.4531
 Epoch 20/50
 7550/7550 20s 3ms/step -
 loss: 2.8463 - mae: 0.4900 - val_loss: 2.4343 - val_mae: 0.4932
 Epoch 21/50
 7550/7550 20s 3ms/step -
 loss: 2.7934 - mae: 0.4841 - val_loss: 2.2943 - val_mae: 0.4580
 Epoch 22/50
 7550/7550 21s 3ms/step -
 loss: 2.7122 - mae: 0.4799 - val_loss: 2.2395 - val_mae: 0.4412
 Epoch 23/50
 7550/7550 21s 3ms/step -
 loss: 2.6786 - mae: 0.4718 - val_loss: 2.2519 - val_mae: 0.4378
 Epoch 24/50
 7550/7550 22s 3ms/step -
 loss: 2.7113 - mae: 0.4678 - val_loss: 2.3175 - val_mae: 0.4527
 Epoch 25/50
 7550/7550 23s 3ms/step -
 loss: 2.6975 - mae: 0.4697 - val_loss: 2.1978 - val_mae: 0.4388
 Epoch 26/50
 7550/7550 21s 3ms/step -
 loss: 2.7644 - mae: 0.4762 - val_loss: 2.5314 - val_mae: 0.5120
 Epoch 27/50
 7550/7550 20s 3ms/step -
 loss: 2.6065 - mae: 0.4600 - val_loss: 2.6210 - val_mae: 0.4839
 Epoch 28/50
 7550/7550 20s 3ms/step -
 loss: 2.6119 - mae: 0.4622 - val_loss: 2.2935 - val_mae: 0.4496
 Epoch 29/50
 7550/7550 19s 3ms/step -
 loss: 2.5467 - mae: 0.4572 - val_loss: 2.2713 - val_mae: 0.4291
 Epoch 30/50
 7550/7550 19s 3ms/step -
 loss: 2.5296 - mae: 0.4528 - val_loss: 2.0601 - val_mae: 0.4250
 Epoch 31/50

7550/7550 19s 3ms/step -
 loss: 2.5385 - mae: 0.4474 - val_loss: 2.2361 - val_mae: 0.4351
 Epoch 32/50
 7550/7550 19s 3ms/step -
 loss: 2.4830 - mae: 0.4436 - val_loss: 2.2456 - val_mae: 0.4683
 Epoch 33/50
 7550/7550 19s 3ms/step -
 loss: 2.5024 - mae: 0.4475 - val_loss: 2.1952 - val_mae: 0.4444
 Epoch 34/50
 7550/7550 20s 3ms/step -
 loss: 2.3724 - mae: 0.4337 - val_loss: 2.2242 - val_mae: 0.4049
 Epoch 35/50
 7550/7550 19s 3ms/step -
 loss: 2.3864 - mae: 0.4283 - val_loss: 2.0833 - val_mae: 0.3840
 Epoch 36/50
 7550/7550 19s 2ms/step -
 loss: 2.3655 - mae: 0.4325 - val_loss: 1.9880 - val_mae: 0.3919
 Epoch 37/50
 7550/7550 20s 3ms/step -
 loss: 2.2759 - mae: 0.4242 - val_loss: 1.9845 - val_mae: 0.3989
 Epoch 38/50
 7550/7550 23s 3ms/step -
 loss: 2.3150 - mae: 0.4261 - val_loss: 1.8645 - val_mae: 0.3617
 Epoch 39/50
 7550/7550 20s 3ms/step -
 loss: 2.3536 - mae: 0.4332 - val_loss: 1.8498 - val_mae: 0.3678
 Epoch 40/50
 7550/7550 19s 3ms/step -
 loss: 2.2434 - mae: 0.4196 - val_loss: 2.0267 - val_mae: 0.4022
 Epoch 41/50
 7550/7550 19s 2ms/step -
 loss: 2.2185 - mae: 0.4168 - val_loss: 1.9775 - val_mae: 0.3958
 Epoch 42/50
 7550/7550 19s 3ms/step -
 loss: 2.2516 - mae: 0.4143 - val_loss: 1.8752 - val_mae: 0.3801
 Epoch 43/50
 7550/7550 20s 3ms/step -
 loss: 2.2185 - mae: 0.4149 - val_loss: 1.7228 - val_mae: 0.3343
 Epoch 44/50
 7550/7550 19s 3ms/step -
 loss: 2.2425 - mae: 0.4182 - val_loss: 1.7470 - val_mae: 0.3686
 Epoch 45/50
 7550/7550 19s 3ms/step -
 loss: 2.1616 - mae: 0.4100 - val_loss: 1.8580 - val_mae: 0.3856
 Epoch 46/50
 7550/7550 19s 2ms/step -
 loss: 2.2155 - mae: 0.4110 - val_loss: 1.8101 - val_mae: 0.3874
 Epoch 47/50


```

7550/7550          19s 3ms/step -
loss: 2.1695 - mae: 0.4067 - val_loss: 1.7064 - val_mae: 0.3533
Epoch 48/50
7550/7550          20s 3ms/step -
loss: 2.1098 - mae: 0.4070 - val_loss: 1.9816 - val_mae: 0.3794
Epoch 49/50
7550/7550          20s 3ms/step -
loss: 2.1072 - mae: 0.4068 - val_loss: 1.7969 - val_mae: 0.3390
Epoch 50/50
7550/7550          18s 2ms/step -
loss: 2.0907 - mae: 0.3986 - val_loss: 1.7812 - val_mae: 0.3507

```

Put Options:

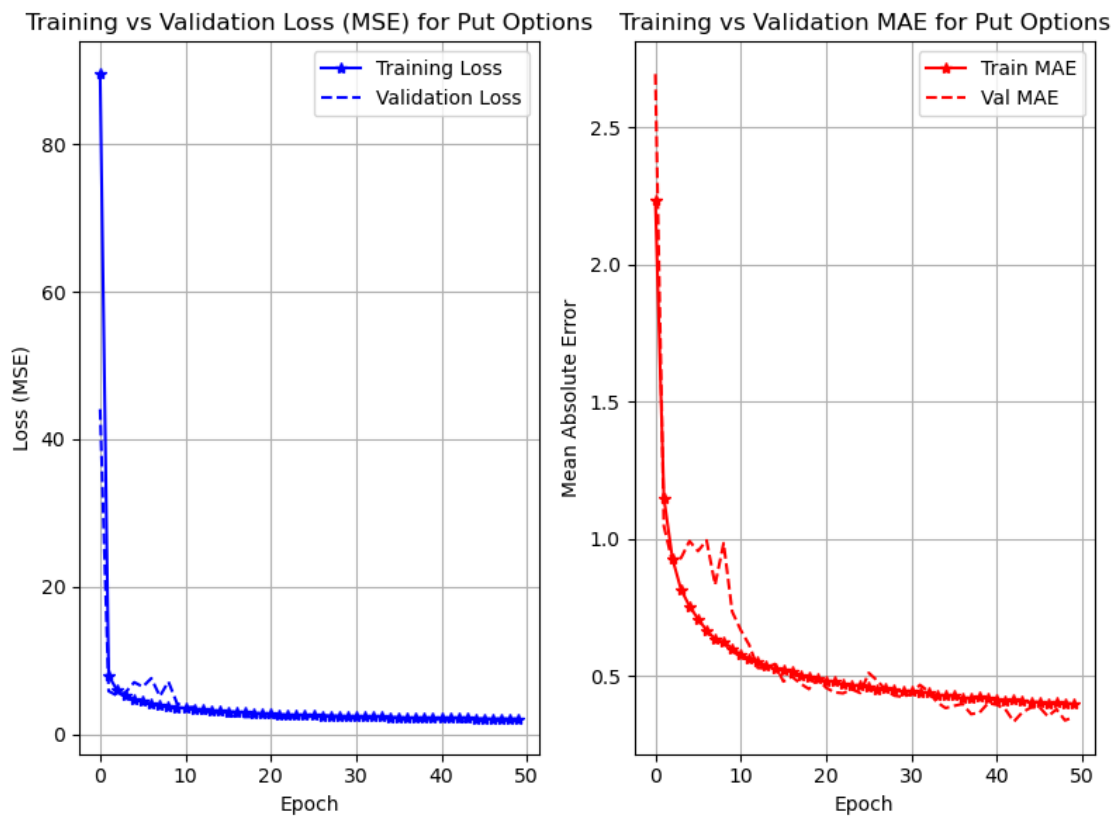
Validation MSE: 1.7063614130020142

Validation MAE: 0.3533361852169037

Test MSE: 1.6476459503173828

Test MAE: 0.35210031270980835

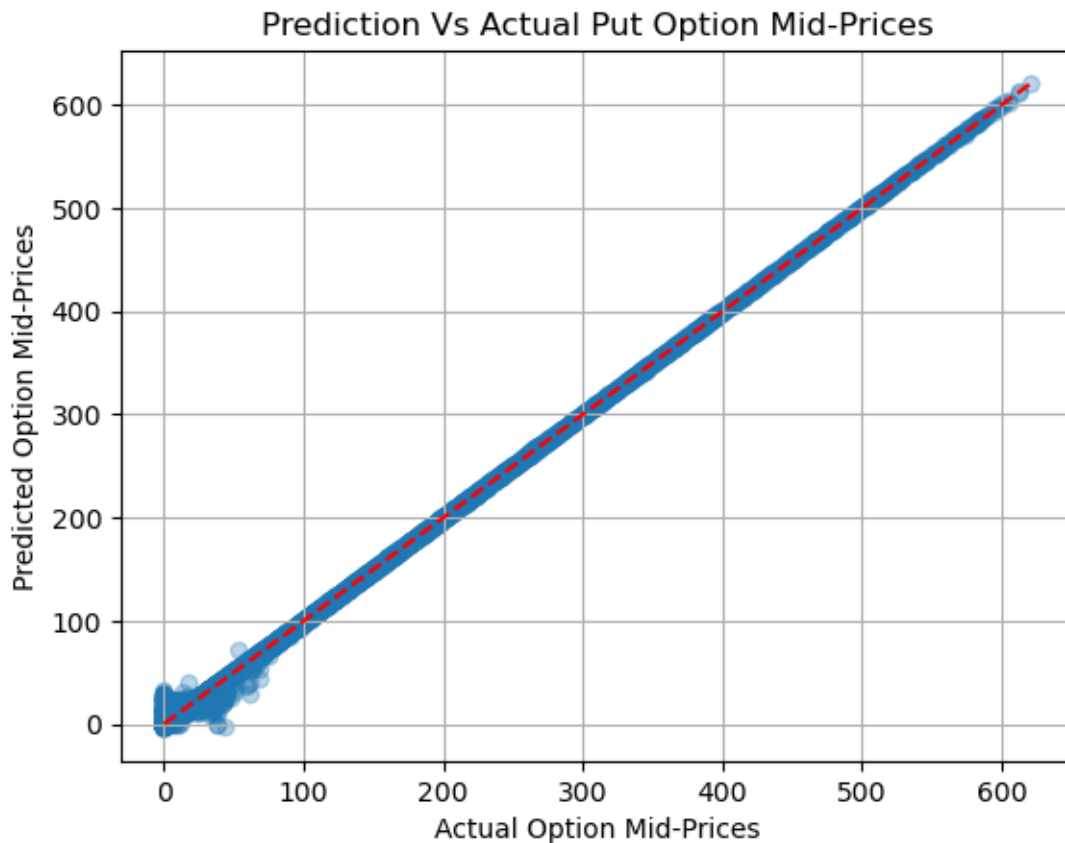
```
[13]: plot_learning_curves(history_puts, "Put")
```



```
[14]: get_prediction_plot(model_puts, X_test_scaled_p, y_test_p, "Put")
```

1619/1619

2s 1ms/step



3.3 Hyperparameter search

To optimize the Neural Network model, we want to identify the best hyperparameter configuration which predicts the option's `mid_price` closest, to do this I will implement a randomized search.

A few of the hyperparameters to optimize are: * Number of Hidden Layers * Number of Neurons per layer * Neuron activation function * Neuron l2 regularization * Learning rate * Optimizer

To find the best configuration, we will select the number of hidden layers between 1 and 5, as adding more may be redundant. For the number of neurons, we will select a value between 10 and 100 neurons. Since deep neural networks often perform better with 'swish' activation rather than with 'ReLU', we will select the activation function between those two at random to assess which performs better. Our l2 regularizer and learning rate will be chosen from a range between $1e-6$ and $1e-2$ to allow a broader range of values in which to assess performance on. Lastly, the optimizer will be selected at random as either the 'Adam' optimizer or the adam optimizer with nesterov momentum ('Nadam') to assess which fits the model better. Once again we will include early stopping to avoid unnecessary training, but this time with a patience of 5, as we will only

be training on 20 epochs instead of 50 due to time constraints. This randomized search will run for 30 iterations in order to assess the optimal parameters, which we will then compare the best model here against the baseline model on their test set performance.

3.3.1 Call Option

```
[28]: def reset_session(seed=42):
    tf.random.set_seed(seed)
    np.random.seed(seed)
    tf.keras.backend.clear_session()

def build_model(hp):
    reset_session()

    model = tf.keras.Sequential()
    model.add(tf.keras.Input(shape = (X_train_c.shape[1],)))

    n_hidden = hp.Int("n_hidden", 1, 5)
    for i in range(n_hidden):
        model.add(tf.keras.layers.Dense(
            hp.Int(f"n_neurons_{i+1}", 10, 100),
            activation = hp.Choice("activation", ["relu", 'swish']),
            kernel_regularizer = tf.keras.regularizers.l2(
                hp.Float("l2", 1e-6, 1e-4, sampling="log")
            ),
            kernel_initializer = "he_normal"
        ))

    model.add(tf.keras.layers.Dense(1))

    learning_rate = hp.Float("learning_rate", 1e-4, 1e-2, sampling = "log")
    optimizer_name = hp.Choice("optimizer", ["Adam", "Nadam"])

    if optimizer_name == "Adam":
        optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    else:
        optimizer = tf.keras.optimizers.Nadam(learning_rate=learning_rate)

    model.compile(loss="mse", metrics = ['mae'], optimizer=optimizer)
    return model

early_stop = EarlyStopping(
    monitor = "val_loss",
    patience = 10,
    mode = "min",
    restore_best_weights = True
)
```

```
[16]: # Train on subsets of training data for speed
X_train_c_sub = X_train_scaled_c[:,5]
X_valid_c_sub = X_valid_scaled_c[:,5]

y_train_c_sub = y_train_c[:,5]
y_valid_c_sub = y_valid_c[:,5]

random_search_tuner = kt.RandomSearch(
    build_model,
    objective = "val_loss",
    seed = 42,
    max_trials = 30,
    overwrite = True
)

random_search_tuner.search(X_train_c_sub, y_train_c_sub, epochs = 30,
                           validation_data = (X_valid_c_sub, y_valid_c_sub), verbose = 1,
                           callbacks=[early_stop])
```

Trial 30 Complete [00h 02m 18s]

val_loss: 0.7884814739227295

Best val_loss So Far: 0.77161705493927

Total elapsed time: 00h 49m 21s

```
[17]: best_hyperparams_c = random_search_tuner.get_best_hyperparameters(num_trials = 1)[0]
best_hyperparams_c.values
```

```
[17]: {'n_hidden': 2,
      'n_neurons_1': 72,
      'activation': 'relu',
      'l2': 4.938970115853091e-06,
      'learning_rate': 0.002220709089529535,
      'optimizer': 'Nadam',
      'n_neurons_2': 82,
      'n_neurons_3': 20,
      'n_neurons_4': 88,
      'n_neurons_5': 72}
```

```
[29]: best_model_c = build_model(best_hyperparams_c)

start = time.time()
history_best_c = best_model_c.fit(X_train_scaled_c, y_train_c,
                                   validation_data = (X_valid_scaled_c, y_valid_c),
                                   epochs = 50, verbose = 1, callbacks=[early_stop])
```

```

end = time.time()
tuned_model_time_c = end - start

test_loss_best_c, test_mae_best_c = best_model_c.evaluate(X_test_scaled_c, y_test_c)
print("\nTest MSE:", test_loss_best_c)
print("Test MAE:", test_mae_best_c)

best_model_c.save("best_model_calls.keras")

```

```

Epoch 1/50
7549/7549          19s 2ms/step -
loss: 297.5094 - mae: 4.8474 - val_loss: 3.8032 - val_mae: 0.8372
Epoch 2/50
7549/7549          17s 2ms/step -
loss: 2.5466 - mae: 0.7101 - val_loss: 1.1366 - val_mae: 0.4629
Epoch 3/50
7549/7549          18s 2ms/step -
loss: 1.5364 - mae: 0.5486 - val_loss: 1.0277 - val_mae: 0.4217
Epoch 4/50
7549/7549          17s 2ms/step -
loss: 1.4010 - mae: 0.4953 - val_loss: 1.1078 - val_mae: 0.4389
Epoch 5/50
7549/7549          17s 2ms/step -
loss: 1.2654 - mae: 0.4526 - val_loss: 1.0340 - val_mae: 0.4177
Epoch 6/50
7549/7549          17s 2ms/step -
loss: 1.2295 - mae: 0.4362 - val_loss: 0.8626 - val_mae: 0.3538
Epoch 7/50
7549/7549          18s 2ms/step -
loss: 1.1839 - mae: 0.4171 - val_loss: 0.9494 - val_mae: 0.3757
Epoch 8/50
7549/7549          17s 2ms/step -
loss: 1.1492 - mae: 0.4038 - val_loss: 0.9188 - val_mae: 0.3727
Epoch 9/50
7549/7549          20s 3ms/step -
loss: 1.1175 - mae: 0.3904 - val_loss: 0.9182 - val_mae: 0.3548
Epoch 10/50
7549/7549          18s 2ms/step -
loss: 1.1023 - mae: 0.3840 - val_loss: 0.8153 - val_mae: 0.3297
Epoch 11/50
7549/7549          18s 2ms/step -
loss: 1.0755 - mae: 0.3736 - val_loss: 0.8254 - val_mae: 0.3321
Epoch 12/50
7549/7549          18s 2ms/step -
loss: 1.0651 - mae: 0.3676 - val_loss: 0.9133 - val_mae: 0.3697
Epoch 13/50

```

7549/7549 18s 2ms/step -
 loss: 1.0614 - mae: 0.3645 - val_loss: 0.7650 - val_mae: 0.2974
 Epoch 14/50
 7549/7549 17s 2ms/step -
 loss: 1.0629 - mae: 0.3615 - val_loss: 0.7468 - val_mae: 0.2962
 Epoch 15/50
 7549/7549 17s 2ms/step -
 loss: 1.0411 - mae: 0.3559 - val_loss: 0.7396 - val_mae: 0.2928
 Epoch 16/50
 7549/7549 18s 2ms/step -
 loss: 1.0239 - mae: 0.3508 - val_loss: 0.7097 - val_mae: 0.2843
 Epoch 17/50
 7549/7549 17s 2ms/step -
 loss: 1.0186 - mae: 0.3479 - val_loss: 0.7171 - val_mae: 0.2946
 Epoch 18/50
 7549/7549 17s 2ms/step -
 loss: 1.0083 - mae: 0.3450 - val_loss: 0.7078 - val_mae: 0.2786
 Epoch 19/50
 7549/7549 17s 2ms/step -
 loss: 0.9959 - mae: 0.3412 - val_loss: 0.7017 - val_mae: 0.2838
 Epoch 20/50
 7549/7549 16s 2ms/step -
 loss: 1.0022 - mae: 0.3411 - val_loss: 0.7240 - val_mae: 0.2940
 Epoch 21/50
 7549/7549 17s 2ms/step -
 loss: 0.9892 - mae: 0.3373 - val_loss: 0.6679 - val_mae: 0.2665
 Epoch 22/50
 7549/7549 17s 2ms/step -
 loss: 0.9856 - mae: 0.3360 - val_loss: 0.6800 - val_mae: 0.2711
 Epoch 23/50
 7549/7549 17s 2ms/step -
 loss: 0.9844 - mae: 0.3354 - val_loss: 0.7091 - val_mae: 0.2962
 Epoch 24/50
 7549/7549 17s 2ms/step -
 loss: 0.9814 - mae: 0.3342 - val_loss: 0.6491 - val_mae: 0.2529
 Epoch 25/50
 7549/7549 16s 2ms/step -
 loss: 0.9783 - mae: 0.3314 - val_loss: 0.7266 - val_mae: 0.3017
 Epoch 26/50
 7549/7549 17s 2ms/step -
 loss: 0.9724 - mae: 0.3294 - val_loss: 0.6871 - val_mae: 0.2839
 Epoch 27/50
 7549/7549 17s 2ms/step -
 loss: 0.9687 - mae: 0.3283 - val_loss: 0.6631 - val_mae: 0.2686
 Epoch 28/50
 7549/7549 17s 2ms/step -
 loss: 0.9567 - mae: 0.3246 - val_loss: 0.7090 - val_mae: 0.2964
 Epoch 29/50

```

7549/7549          16s 2ms/step -
loss: 0.9532 - mae: 0.3231 - val_loss: 0.6870 - val_mae: 0.2806
Epoch 30/50
7549/7549          17s 2ms/step -
loss: 0.9469 - mae: 0.3225 - val_loss: 0.7291 - val_mae: 0.3064
Epoch 31/50
7549/7549          17s 2ms/step -
loss: 0.9513 - mae: 0.3227 - val_loss: 0.7262 - val_mae: 0.3052
Epoch 32/50
7549/7549          18s 2ms/step -
loss: 0.9395 - mae: 0.3198 - val_loss: 0.6934 - val_mae: 0.2924
Epoch 33/50
7549/7549          17s 2ms/step -
loss: 0.9350 - mae: 0.3172 - val_loss: 0.7055 - val_mae: 0.2865
Epoch 34/50
7549/7549          17s 2ms/step -
loss: 0.9326 - mae: 0.3160 - val_loss: 0.7484 - val_mae: 0.3125
1617/1617          3s 2ms/step -
loss: 0.7468 - mae: 0.2536

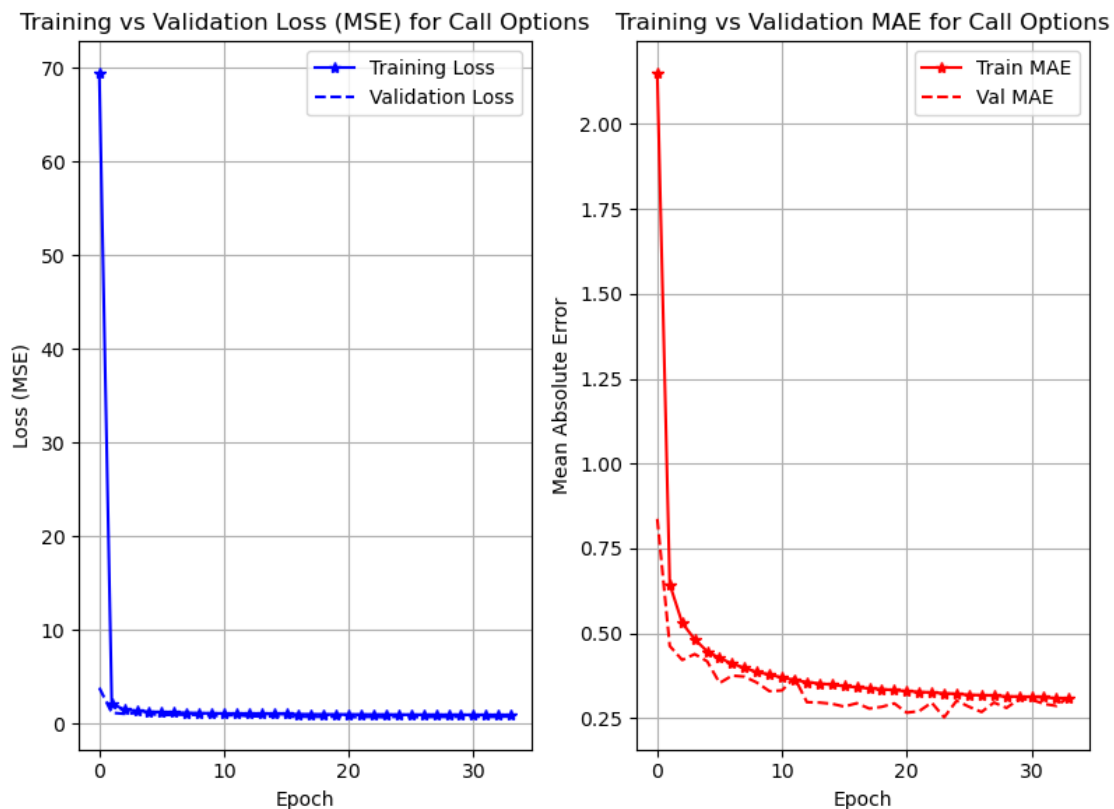
```

```

Test MSE: 0.844775915145874
Test MAE: 0.2551146447658539

```

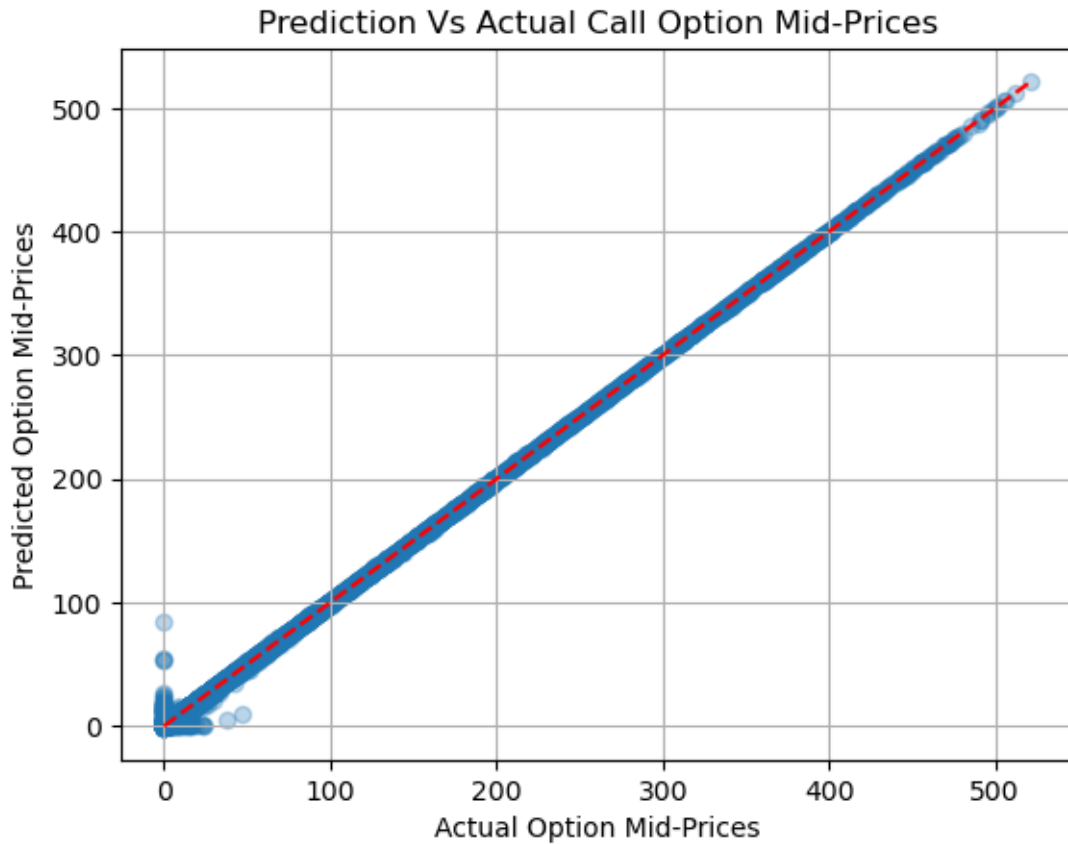
```
[30]: plot_learning_curves(history_best_c, "Call")
```



```
[31]: get_prediction_plot(best_model_c, X_test_scaled_c,  
                          y_test_c, "Call")
```

1617/1617

2s 1ms/step



3.3.2 Put Option

```
[21]: X_train_p_sub = X_train_scaled_p[:,5]  
X_valid_p_sub = X_valid_scaled_p[:,5]  
  
y_train_p_sub = y_train_p[:,5]  
y_valid_p_sub = y_valid_p[:,5]  
  
random_search_tuner = kt.RandomSearch(  
    build_model,  
    objective = "val_loss",  
    seed = 42,
```



```

        max_trials = 30,
        overwrite = True
    )

    random_search_tuner.search(X_train_p_sub, y_train_p_sub, epochs = 30,
                              validation_data = (X_valid_p_sub, y_valid_p_sub), verbose = 1,
                              callbacks=[early_stop])

```

Trial 30 Complete [00h 02m 29s]

val_loss: 3.373318910598755

Best val_loss So Far: 2.8543083667755127

Total elapsed time: 00h 55m 25s

```

[22]: best_hyperparams_p = random_search_tuner.get_best_hyperparameters(num_trials = 1)[0]
      best_hyperparams_p.values

```

```

[22]: {'n_hidden': 5,
      'n_neurons_1': 100,
      'activation': 'relu',
      'l2': 1.4513514342725857e-05,
      'learning_rate': 0.0009452767099282696,
      'optimizer': 'Nadam',
      'n_neurons_2': 72,
      'n_neurons_3': 45,
      'n_neurons_4': 21,
      'n_neurons_5': 57}

```

```

[32]: best_model_p = build_model(best_hyperparams_p)

start = time.time()
history_best_p = best_model_p.fit(X_train_scaled_p, y_train_p,
                                  validation_data = (X_valid_scaled_p, y_valid_p),
                                  epochs = 50, verbose = 1, callbacks=[early_stop])
end = time.time()
tuned_model_time_p = end - start

test_loss_best_p, test_mae_best_p = best_model_p.evaluate(X_test_scaled_p, y_test_p)
print("\nTest MSE:", test_loss_best_p)
print("Test MAE:", test_mae_best_p)

best_model_p.save("best_model_puts.keras")

```

Epoch 1/50

7550/7550 22s 2ms/step -
 loss: 318.0915 - mae: 3.6989 - val_loss: 10.8060 - val_mae: 1.3758
 Epoch 2/50
 7550/7550 20s 3ms/step -
 loss: 11.0536 - mae: 1.2549 - val_loss: 7.2598 - val_mae: 1.1116
 Epoch 3/50
 7550/7550 23s 3ms/step -
 loss: 7.9689 - mae: 1.0529 - val_loss: 3.9439 - val_mae: 0.7532
 Epoch 4/50
 7550/7550 20s 3ms/step -
 loss: 6.5212 - mae: 0.9298 - val_loss: 13.1403 - val_mae: 1.3690
 Epoch 5/50
 7550/7550 19s 3ms/step -
 loss: 5.5278 - mae: 0.8274 - val_loss: 4.5645 - val_mae: 0.8259
 Epoch 6/50
 7550/7550 19s 3ms/step -
 loss: 4.7953 - mae: 0.7639 - val_loss: 3.8391 - val_mae: 0.7064
 Epoch 7/50
 7550/7550 20s 3ms/step -
 loss: 4.4247 - mae: 0.7106 - val_loss: 5.8779 - val_mae: 0.9263
 Epoch 8/50
 7550/7550 20s 3ms/step -
 loss: 4.1753 - mae: 0.6839 - val_loss: 4.7803 - val_mae: 0.8759
 Epoch 9/50
 7550/7550 19s 3ms/step -
 loss: 4.0149 - mae: 0.6506 - val_loss: 3.0456 - val_mae: 0.5631
 Epoch 10/50
 7550/7550 19s 2ms/step -
 loss: 3.7992 - mae: 0.6221 - val_loss: 2.7964 - val_mae: 0.5587
 Epoch 11/50
 7550/7550 19s 2ms/step -
 loss: 3.5853 - mae: 0.6012 - val_loss: 2.5173 - val_mae: 0.4892
 Epoch 12/50
 7550/7550 19s 2ms/step -
 loss: 3.4402 - mae: 0.5796 - val_loss: 2.8135 - val_mae: 0.4731
 Epoch 13/50
 7550/7550 20s 3ms/step -
 loss: 3.3973 - mae: 0.5604 - val_loss: 2.5415 - val_mae: 0.5338
 Epoch 14/50
 7550/7550 20s 3ms/step -
 loss: 3.4829 - mae: 0.5665 - val_loss: 4.8309 - val_mae: 0.7896
 Epoch 15/50
 7550/7550 20s 3ms/step -
 loss: 3.3990 - mae: 0.5645 - val_loss: 2.8756 - val_mae: 0.4989
 Epoch 16/50
 7550/7550 19s 2ms/step -
 loss: 3.1936 - mae: 0.5405 - val_loss: 2.8150 - val_mae: 0.4970
 Epoch 17/50

7550/7550 18s 2ms/step -
 loss: 3.1361 - mae: 0.5273 - val_loss: 3.2202 - val_mae: 0.5897
 Epoch 18/50
 7550/7550 19s 2ms/step -
 loss: 3.1547 - mae: 0.5292 - val_loss: 2.1182 - val_mae: 0.3878
 Epoch 19/50
 7550/7550 19s 3ms/step -
 loss: 2.8985 - mae: 0.5019 - val_loss: 2.4412 - val_mae: 0.4894
 Epoch 20/50
 7550/7550 19s 3ms/step -
 loss: 2.9947 - mae: 0.4952 - val_loss: 2.0268 - val_mae: 0.4126
 Epoch 21/50
 7550/7550 20s 3ms/step -
 loss: 2.9918 - mae: 0.5001 - val_loss: 2.8676 - val_mae: 0.5401
 Epoch 22/50
 7550/7550 19s 3ms/step -
 loss: 2.9596 - mae: 0.4990 - val_loss: 2.2578 - val_mae: 0.4049
 Epoch 23/50
 7550/7550 19s 3ms/step -
 loss: 2.8667 - mae: 0.4750 - val_loss: 2.3602 - val_mae: 0.4485
 Epoch 24/50
 7550/7550 19s 2ms/step -
 loss: 2.8646 - mae: 0.4764 - val_loss: 2.0667 - val_mae: 0.3668
 Epoch 25/50
 7550/7550 18s 2ms/step -
 loss: 2.9535 - mae: 0.4819 - val_loss: 2.4650 - val_mae: 0.4817
 Epoch 26/50
 7550/7550 18s 2ms/step -
 loss: 2.8617 - mae: 0.4824 - val_loss: 2.6603 - val_mae: 0.4992
 Epoch 27/50
 7550/7550 18s 2ms/step -
 loss: 2.9167 - mae: 0.4800 - val_loss: 2.2304 - val_mae: 0.3618
 Epoch 28/50
 7550/7550 23s 3ms/step -
 loss: 2.6285 - mae: 0.4590 - val_loss: 1.9010 - val_mae: 0.3689
 Epoch 29/50
 7550/7550 19s 2ms/step -
 loss: 2.6563 - mae: 0.4557 - val_loss: 2.1501 - val_mae: 0.3998
 Epoch 30/50
 7550/7550 19s 3ms/step -
 loss: 2.9138 - mae: 0.4717 - val_loss: 2.0838 - val_mae: 0.4328
 Epoch 31/50
 7550/7550 19s 2ms/step -
 loss: 2.5626 - mae: 0.4484 - val_loss: 1.9963 - val_mae: 0.3620
 Epoch 32/50
 7550/7550 20s 3ms/step -
 loss: 2.5251 - mae: 0.4394 - val_loss: 2.1065 - val_mae: 0.4008
 Epoch 33/50

7550/7550 20s 3ms/step -
loss: 2.5577 - mae: 0.4456 - val_loss: 1.8781 - val_mae: 0.3679
Epoch 34/50

7550/7550 19s 3ms/step -
loss: 2.4690 - mae: 0.4381 - val_loss: 2.0685 - val_mae: 0.4209
Epoch 35/50

7550/7550 20s 3ms/step -
loss: 2.4860 - mae: 0.4380 - val_loss: 2.0144 - val_mae: 0.3624
Epoch 36/50

7550/7550 20s 3ms/step -
loss: 2.5600 - mae: 0.4327 - val_loss: 1.8152 - val_mae: 0.3553
Epoch 37/50

7550/7550 19s 2ms/step -
loss: 2.4994 - mae: 0.4256 - val_loss: 1.9254 - val_mae: 0.3468
Epoch 38/50

7550/7550 19s 3ms/step -
loss: 2.4538 - mae: 0.4296 - val_loss: 1.8603 - val_mae: 0.3864
Epoch 39/50

7550/7550 19s 2ms/step -
loss: 2.4073 - mae: 0.4283 - val_loss: 1.8291 - val_mae: 0.3703
Epoch 40/50

7550/7550 19s 2ms/step -
loss: 2.4066 - mae: 0.4166 - val_loss: 1.7332 - val_mae: 0.3274
Epoch 41/50

7550/7550 18s 2ms/step -
loss: 2.3233 - mae: 0.4189 - val_loss: 1.7921 - val_mae: 0.3617
Epoch 42/50

7550/7550 18s 2ms/step -
loss: 2.2780 - mae: 0.4164 - val_loss: 1.7321 - val_mae: 0.3151
Epoch 43/50

7550/7550 18s 2ms/step -
loss: 2.3089 - mae: 0.4157 - val_loss: 1.7851 - val_mae: 0.3167
Epoch 44/50

7550/7550 18s 2ms/step -
loss: 2.2808 - mae: 0.4117 - val_loss: 1.7097 - val_mae: 0.3144
Epoch 45/50

7550/7550 18s 2ms/step -
loss: 2.2843 - mae: 0.4067 - val_loss: 1.8567 - val_mae: 0.3805
Epoch 46/50

7550/7550 18s 2ms/step -
loss: 2.1762 - mae: 0.4015 - val_loss: 1.8308 - val_mae: 0.3973
Epoch 47/50

7550/7550 18s 2ms/step -
loss: 2.2376 - mae: 0.4047 - val_loss: 1.6673 - val_mae: 0.3171
Epoch 48/50

7550/7550 18s 2ms/step -
loss: 2.1754 - mae: 0.3979 - val_loss: 1.7649 - val_mae: 0.3192
Epoch 49/50

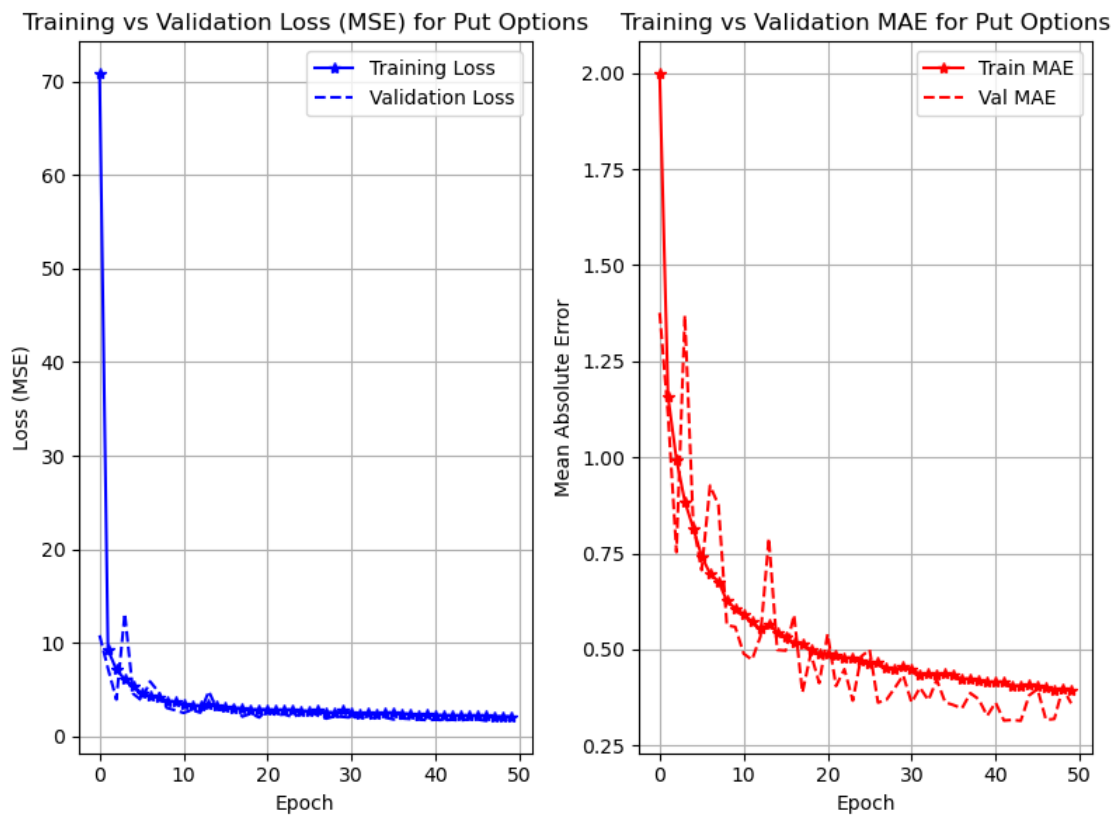
```

7550/7550          18s 2ms/step -
loss: 2.1808 - mae: 0.3970 - val_loss: 1.8743 - val_mae: 0.3997
Epoch 50/50
7550/7550          18s 2ms/step -
loss: 2.2774 - mae: 0.4038 - val_loss: 1.7478 - val_mae: 0.3598
1619/1619          3s 2ms/step -
loss: 1.6900 - mae: 0.3239

Test MSE: 1.654191255569458
Test MAE: 0.3205004334449768

```

```
[33]: plot_learning_curves(history_best_p, "Put")
```

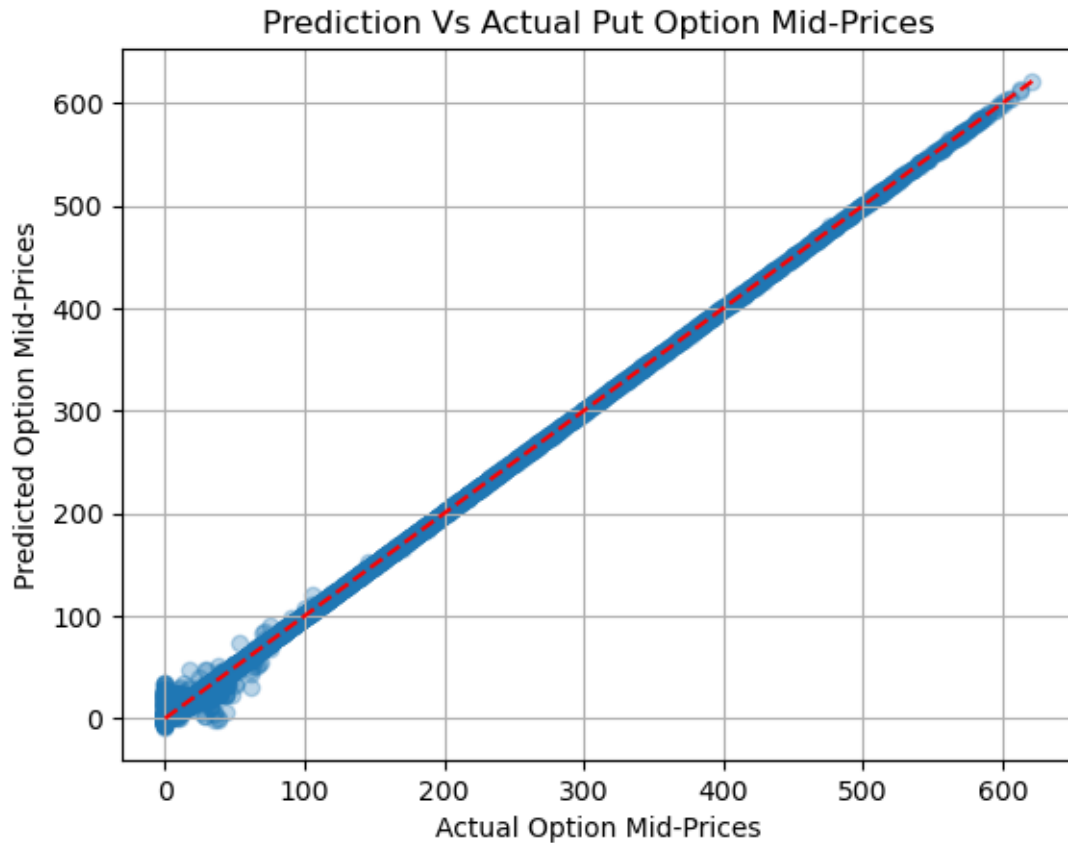


```
[34]: get_prediction_plot(best_model_p, X_test_scaled_p,
                           y_test_p, "Put")
```

```

1619/1619          2s 1ms/step

```



```
[35]: results_df = pd.DataFrame([
    {
        'Model': 'Baseline',
        'Option Type': 'Call Option',
        'Layers': 4,
        'Activation': "ReLU",
        'Optimizer': "Nadam",
        'Epochs': 50,
        'MAE': round(test_mae_c, 4),
        'MSE': round(test_loss_c, 4),
        'Train Time (s)': round(baseline_time_c, 2)
    },
    {
        'Model': 'Tuned',
        'Option Type': 'Call Option',
        'Layers': str(best_hyperparams_c['n_hidden']),
        'Activation': best_hyperparams_c['activation'],
        'Optimizer': best_hyperparams_c['optimizer'],
        'Epochs': max(history_best_c.epoch),
        'MAE': round(test_mae_best_c, 4),
    }
])
```

```

        'MSE': round(test_loss_best_c, 4),
        'Train Time (s)': round(tuned_model_time_c, 2)
    },
    {
        'Model': 'Baseline',
        'Option Type': 'Put Option',
        'Layers': 4,
        'Activation': "ReLU",
        'Optimizer': "Nadam",
        'Epochs': 50,
        'MAE': round(test_mae_p, 4),
        'MSE': round(test_loss_p, 4),
        'Train Time (s)': round(baseline_time_p, 2)
    },
    {
        'Model': 'Tuned',
        'Option Type': 'Put Option',
        'Layers': str(best_hyperparams_p['n_hidden']),
        'Activation': best_hyperparams_p['activation'],
        'Optimizer': best_hyperparams_p['optimizer'],
        'Epochs': max(history_best_p.epoch),
        'MAE': round(test_mae_best_p, 4),
        'MSE': round(test_loss_best_p, 4),
        'Train Time (s)': round(tuned_model_time_p, 2)
    }
])

display(results_df)

```

	Model	Option Type	Layers	Activation	Optimizer	Epochs	MAE	MSE	\
0	Baseline	Call Option	4	ReLU	Nadam	50	0.2697	0.8442	
1	Tuned	Call Option	2	relu	Nadam	33	0.2551	0.8448	
2	Baseline	Put Option	4	ReLU	Nadam	50	0.3521	1.6476	
3	Tuned	Put Option	5	relu	Nadam	49	0.3205	1.6542	

	Train Time (s)
0	667.01
1	587.61
2	981.35
3	955.78